Name – Rishabh Jain || CSP554 || A20495530

Assignment 12

Exercise 1) (4 points)
Read the article "A Big Data Modeling Methodology for Apache Cassandra" available on the blackboard in the 'Articles' section. Provide a ½ page summary including your comments and impressions.

Solution)

Summary
The paper covers traditional data modeling, Cassandra data modeling, conceptual and logical data modeling and application workflow, query driven mapping from a conceptual to a logical data model, and physical data modeling.

Cassandra Data Model:
A CQL table can be conceived of as a collection of divisions that include rows with similar structures. A partition key is unique to each partition in a table, and a clustering key is unique to each row within a partition. A primary key is a combination of a partition key and a clustering key that allows a database row to be uniquely identified. A table schema is a set of columns that includes a primary key. The data type for each column is either primitive (int, text, etc.), complex (set, list, or map), or counter.

CQL, which has a SQL-like syntax, is used to express queries over tables. CQL does not support binary operations such as joins and has a set of query predicates rules that ensure efficiency and scalability.

Conceptual data modelling and application workflow
Understanding the data to be maintained and how a data-driven application needs to access it is required when designing a Cassandra database schema. The ER diagram depicts the former. Application workflow diagrams, which define data access patterns for application tasks, capture the latter.

Query driven mapping

Data Modeling Principles: The four data modeling principles listed below serve as a foundation for translating conceptual data models into logical data models.
DMP1 (Know your data): Understanding the data, which is recorded using a conceptual data model, is the first step in successful database design.
DMP2 (Know your Questions): The second key to a successful database design is knowing your queries, which are captured by an application process.
DMP3 (Data Nesting): Data nesting is the third key to a successful database design.
DMP4 (Data Duplication): Data duplication is the fourth key to a successful database design.

Mapping Rule: - Five mapping rules that facilitate a query-driven move from a conceptual data model to a logical data model are listed below.
MR1 -> Entity and relationship types map to tables, while entities and relationships map to table rows in MR1 (Entities and Relationships).
MR2 -> (Equality Search Attributes): Equality search attributes map to the prefix columns of a table primary key in a query predicate.
MR3 -> (Inequality Search Attributes): A table clustering key column maps to an inequality search attribute utilized in a query predicate.

Name – Rishabh Jain || CSP554 || A20495530

MR4 -> (Ordering Attributes): Ordering attributes, which are supplied in a query, map to clustering key columns in the query's chosen ascending or descending clustering order.

MR5 -> (Key Attributes): Primary key columns are mapped to key attribute types.

Mapping Patterns: Mapping Patterns serve as the basis for automating Cassandra database schema design.

Physical Data Modeling

The final step is the analysis and optimization of a logical data model to produce a physical data model.

Exercise 2) (3 points)

a)

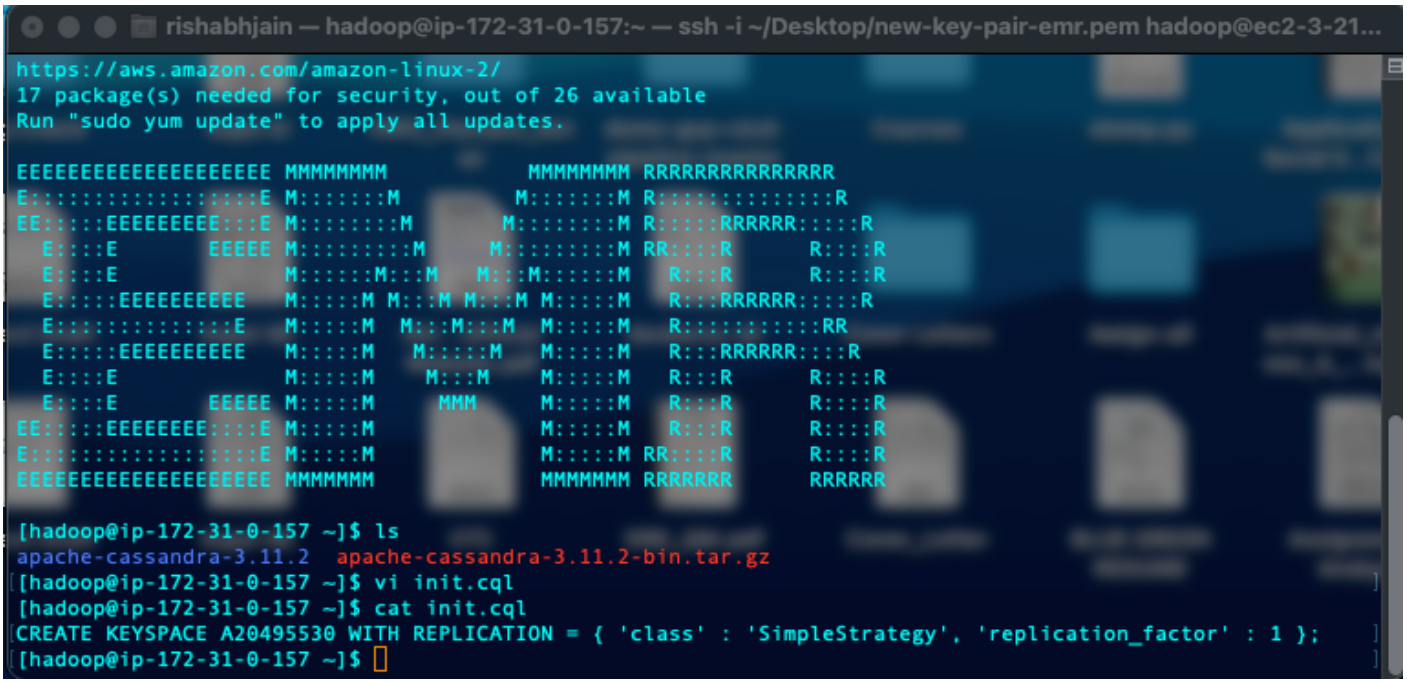

b)

source ‘./init.cql’;

c)

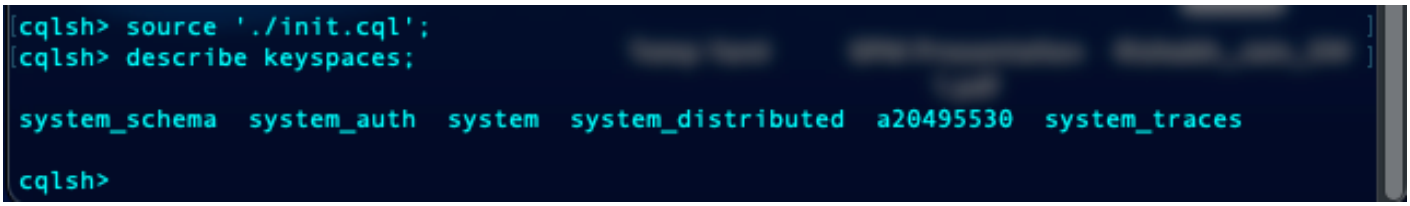

d)

Name – Rishabh Jain || CSP554 || A20495530

```
[hadoop@ip-172-31-0-157 ~]$ ls
apache-cassandra-3.11.2  apache-cassandra-3.11.2-bin.tar.gz  ex2.cql  init.cql
[hadoop@ip-172-31-0-157 ~]$
```

```
cqlsh:a20495530> source './ex2.cql'
cqlsh:a20495530> DESCRIBE TABLE Music;

CREATE TABLE a20495530.music (
    artistname text,
    albumname text,
    cost int,
    numbersold int,
    PRIMARY KEY (artistname, albumname)
) WITH CLUSTERING ORDER BY (albumname ASC)
    AND bloom_filter_fp_chance = 0.01
    AND caching = {'keys': 'ALL', 'rows_per_partition': 'NONE'}
    AND comment = ''
    AND compaction = {'class': 'org.apache.cassandra.db.compaction.SizeTieredCompactionStrategy', 'max_threshold': '32', 'min_threshold': '4'
}
    AND compression = {'chunk_length_in_kb': '64', 'class': 'org.apache.cassandra.io.compress.LZ4Compressor'}
    AND crc_check_chance = 1.0
    AND dclocal_read_repair_chance = 0.1
    AND default_time_to_live = 0
    AND gc_grace_seconds = 864000
    AND max_index_interval = 2048
    AND memtable_flush_period_in_ms = 0
    AND min_index_interval = 128
    AND read_repair_chance = 0.0
    AND speculative_retry = '99PERCENTILE';

cqlsh:a20495530>
```

Exercise 3) (3 points)

a)

```
rishabhjain — hadoop@ip-172-31-0-157:~ — ssh -i ~/Desktop/new-key-pair-emr.pem hadoop@ec2-3-21...
[hadoop@ip-172-31-0-157 ~]$ vi ex3.cql
[hadoop@ip-172-31-0-157 ~]$ cat ex3.cql
insert into Music (artistName, albumName, numberSold, cost)
values ('Mozart', 'Greatest Hits', 100000, 10);

insert into Music (artistName, albumName, numberSold, cost)
values ('Taylor Swift', 'Fearless', 2300000, 15);

insert into Music (artistName, albumName, numberSold, cost)
values ('Black Sabbath', 'Paranoid', 534000, 12);

insert into Music (artistName, albumName, numberSold, cost)
values ('Katy Perry', 'Prism', 800000, 16);

insert into Music (artistName, albumName, numberSold, cost)
values ('Katy Perry', 'Teenage Dream', 750000, 14);

[hadoop@ip-172-31-0-157 ~]$
```

b)

```
cqlsh:a20495530> source './ex3.cql'
cqlsh:a20495530> SELECT * FROM Music;

 artistname    | albumname     | cost | numbersold
---------------+---------------+------+------------
       Mozart | Greatest Hits |   10 |     100000
 Black Sabbath |      Paranoid |   12 |     534000
  Taylor Swift |      Fearless |   15 |    2300000
    Katy Perry |         Prism |   16 |     800000
    Katy Perry | Teenage Dream |   14 |     750000

(5 rows)
cqlsh:a20495530>
```

Name – Rishabh Jain || CSP554 || A20495530

Exercise 4) (2 points)

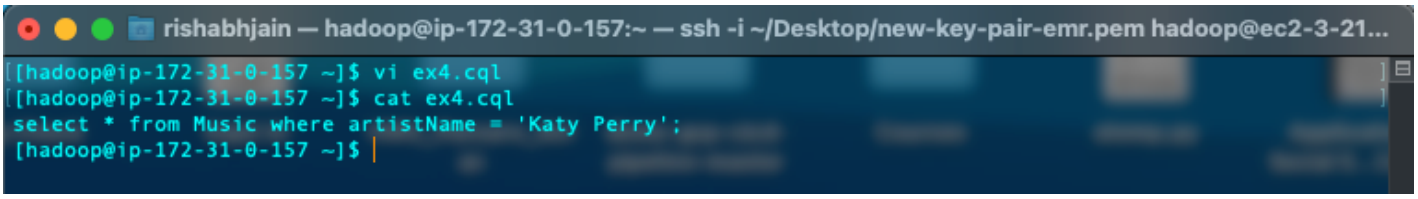

Exercise 5) (2 points)